

LinguAI — What I Need Signed Off Before Launch

Zain, this is the list I want us to go through together before we call v1 done for UmmahTech Innovations. Nothing on here gets checked off just because it's said to be done — I want to see it working with my own eyes.

Three levels, so we're both clear on what's urgent:

- **P0** — Blocks submission to Play Store. Not negotiable.
 - **P1** — Has to be closed within 2 weeks of going live, before we push any marketing.
 - **P2** — Fine to push back, but it needs to be a decision we made on purpose, not something that got forgotten.
-

1. Core App Functionality

- **[P0]** The full learning loop actually works, start to finish — browse lessons, open one, go through flashcards with audio, take the quiz, see results, land back in the review queue. No crashes, no dead ends.
 - **[P0]** Every must-have feature we agreed on is built and demoable — not "it's mostly there," actually working in front of me.
 - **[P0]** Get someone other than yourself to test a feature before it's marked done — even a quick second pair of eyes counts.
 - **[P0]** There's enough real content live at launch — minimum 15 lessons, 150 vocabulary words, and every word has working audio. An app that looks empty on day one kills first impressions.
-

2. Login, Accounts & Abuse Protection

- **[P0]** New signups have to verify their email before they get full access. Right now anyone can register with an email that isn't even theirs.
- **[P0]** Registration, login, and password reset are protected against bots — CAPTCHA or something like Cloudflare Turnstile. Otherwise we'll get scripted fake accounts burning through our AI quota before real users even show up.
- **[P0]** There's an age check at signup, with a clear policy for what happens if someone says they're under 13. This isn't optional — it's a Play Store requirement, not just a nice-to-have.
- **[P1]** Refresh tokens are tied to the device, and there's a "log out of all other devices" option in settings. Right now a stolen token works from anywhere for 7 days.
- **[P1]** Repeated failed logins lock the account temporarily — this should already be part of the plan, just confirm it's actually wired up.

3. The AI Tutor — Safety & Cost

- **[P0]** I need a dashboard, even a basic one, showing how much of our Groq/Gemini free tier we're using in real time. We're building our whole cost model on staying under those limits — flying blind on this is how the AI feature quietly dies mid-launch.
- **[P0]** Before we hit the free-tier ceiling, the system should warn me and fall back gracefully — not just stop working and leave users staring at errors.
- **[P1]** Whatever a user types to the AI tutor gets checked before it's sent off — basic protection against people trying to jailbreak or manipulate it.
- **[P1]** AI responses get screened before they reach the user. This is a language-learning app — kids will use it.
- **[P1]** Pull 50+ real AI responses across different skill levels and actually read them for accuracy. Don't just trust that RAG is working because it's supposed to.

- **[P1]** Test what a brand-new user with zero history sees from the AI — make sure there's no broken placeholder text where their history would normally go.

4. Servers, Uptime & What Happens When Things Break

- **[P0]** Run a real load test at the concurrency we're targeting, on the actual server spec we're paying for. Don't assume it holds up — measure it.
- **[P0]** Get error tracking wired in (Sentry or similar) so I find out about crashes from a dashboard, not from a bad review.
- **[P0]** Basic uptime monitoring pinging the health endpoint, with alerts turned on.
- **[P1]** A staging environment exists — deploys go there first, not straight to production from a commit.
- **[P1]** Actually restore a backup once. Time it. Write down the steps. "We take backups" means nothing if we've never tested pulling one back.
- **[P1]** Simulate an AI provider going down and confirm the fallback logic actually kicks in like it's supposed to.
- **[P1]** A one-page doc: what to check and who does what when the app goes down at 2am.

5. Privacy, Data & Legal

- **[P0]** Privacy policy and terms get reviewed by someone other than whoever wrote the code, before we submit to Play Store.
- **[P0]** The Google Play Data Safety form actually matches what the app really collects and shares — not a generic template.
- **[P0]** Account deletion actually deletes the data within the promised window — test this on a real throwaway account, don't just trust the code.
- **[P1]** We need an answer on the Wiktionary content licensing question — even a single paid one-hour consult with a lawyer is

enough for now. Right now this is just sitting as an unresolved risk, and that's not okay for a paid product.

- **[P1]** Data export actually works and delivers a complete file, tested on a real account.

6. Payments & Subscriptions

- **[P0]** Walk through a real purchase in the Play Store sandbox — subscribe, webhook fires, access unlocks, no app restart needed.
- **[P0]** Send the same payment webhook twice on purpose and confirm it doesn't double-charge logic or corrupt the subscription state.
- **[P0]** Subscription status gets re-checked against the backend every time the app opens — never trusted purely from local cache.
- **[P1]** Test "restore purchases" on a second device with the same account.

7. Testing

- **[P0]** Backend test coverage hits at least 70%, and this is enforced automatically — the build should fail if it drops below that.
- **[P0]** Run a security scan (OWASP ZAP or equivalent) and get it to zero high-severity findings.
- **[P1]** At minimum, the login flow, the purchase flow, and the offline sync logic have automated tests on the app side too — these are the three places where a silent bug would hurt the most.
- **[P1]** Someone spends an afternoon actually trying to break the login and payment webhook endpoints — replay attempts, basic injection, auth bypass.

8. Language & Accessibility

- **[P0]** Switching between English and Urdu works from settings without restarting the app.
- **[P0]** Urdu mode actually mirrors correctly — layout, icons, text alignment, all of it.
- **[P0]** The Urdu font renders properly on a real mid-range phone, not just the simulator.
- **[P2]** We consciously decide whether screen-reader and font-scaling support is in v1 or pushed to v2 — I just want that to be a decision, not something nobody thought about.

9. Speed & Performance

- **[P0]** Cold start under 3 seconds on a mid-range phone, not a flagship.
- **[P0]** Flashcard animations stay smooth — no visible stutter.
- **[P0]** Regular API calls (not AI ones) respond in under half a second under normal load.
- **[P1]** Final app size stays under 50MB.

10. Ready to Actually Launch

- **[P0]** Push to main triggers a full build-and-deploy pipeline automatically, and I've watched it run at least once end-to-end.
 - **[P0]** There's a real support email or channel users can reach — not a placeholder.
 - **[P1]** If a deploy breaks something, there's a documented way to roll it back fast.
 - **[P1]** I get a handoff doc — environment variables, deployment steps, how the pieces fit together — so this doesn't fall apart if the dev is unreachable for two weeks.
-

How we'll use this

Let's go through this list item by item, live. P0 items are the gate — nothing goes to Play Store until every one of them is checked off in front of me. For P1 and P2 items, I want an actual date attached, not "we'll get to it."